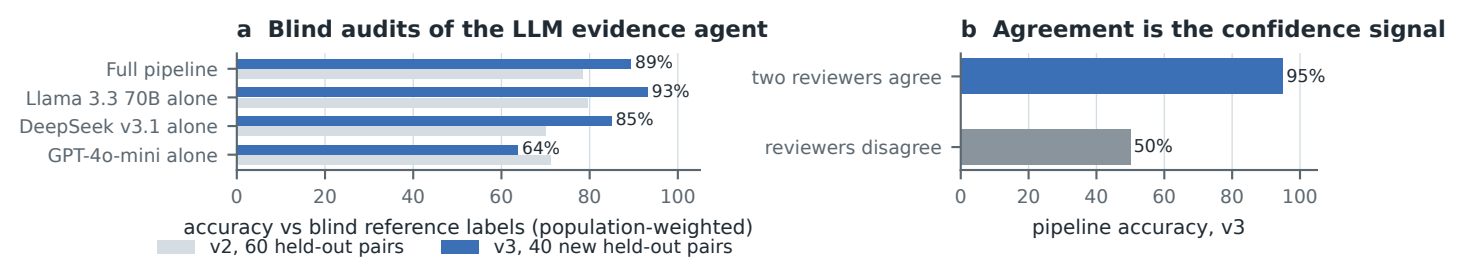
a Blind audits of the LLM evidence agent
Full pipeline
Llama 3.3 70B alone
DeepSeek v3.1 alone
GPT-4o-mini alone
89%
93%
85%
64%
0
20
40
60
80
100
accuracy vs blind reference labels (population-weighted)
v2, 60 held-out pairs
v3, 40 new held-out pairs
b Agreement is the confidence signal
two reviewers agree
reviewers disagree
95%
50%
0
20
40
60
80
100
pipeline accuracy, v3